

Project_03_Sentiment Analysis IMDB Dataset

<https://www.kaggle.com/code/lakshmi25npathi/sentiment-analysis-of-imdb-movie-reviews>
(<https://www.kaggle.com/code/lakshmi25npathi/sentiment-analysis-of-imdb-movie-reviews>)

```
In [1]: 1 import numpy as np
2 import pandas as pd
3 import nltk
4 from nltk.sentiment.vader import SentimentIntensityAnalyzer
5 import re
6 from textblob import TextBlob
7 from wordcloud import WordCloud
8 import seaborn as sns
9 import matplotlib.pyplot as plt
10 import cufflinks as cf
11 %matplotlib inline
12 from plotly.offline import init_notebook_mode, iplot
13 init_notebook_mode(connected = True)
14 cf.go_offline();
15 import plotly.graph_objs as go
16 from plotly.subplots import make_subplots
17
18 import warnings
19 warnings.filterwarnings('ignore')
20 warnings.warn('this will not show')
21
22 pd.set_option('display.max_columns', None)
23
24
25 from nltk.corpus import stopwords
26 from nltk.stem import SnowballStemmer
27 from sklearn.feature_extraction.text import CountVectorizer
28
29 from collections import Counter
30 from numpy import where
31
32 from imblearn.over_sampling import SMOTE
33 from sklearn.decomposition import PCA
34
35 from sklearn.preprocessing import OneHotEncoder
36 from sklearn.preprocessing import StandardScaler
37
38 from sklearn.model_selection import train_test_split
39 from sklearn.linear_model import LogisticRegression
40 from sklearn.metrics import accuracy_score
41 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
42 from sklearn import metrics
43
44 from scipy.sparse import hstack, vstack
45
46 from prettytable import PrettyTable
47 from scipy.stats import loguniform # Log-uniform is useful for searching
48 from sklearn.model_selection import RepeatedStratifiedKFold, RandomizedSe
```

C:\ProgramData\anaconda3\Lib\site-packages\paramiko\transport.py:219: CryptographyDeprecationWarning:

Blowfish has been deprecated

In [2]: 1 nltk.download('stopwords')

[nltk_data] Downloading package stopwords to
 [nltk_data] C:\Users\blues\AppData\Roaming\nltk_data...
 [nltk_data] Package stopwords is already up-to-date!

Out[2]: True

In [3]: 1 nltk.download('punkt')

[nltk_data] Downloading package punkt to
 [nltk_data] C:\Users\blues\AppData\Roaming\nltk_data...
 [nltk_data] Package punkt is already up-to-date!

Out[3]: True

In [4]: 1 df = pd.read_csv("IMDB Dataset.csv", delimiter=',', engine='python', enco
 2 df.head())

Out[4]:

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production. The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive

In [5]: 1 df.shape

Out[5]: (50000, 2)

In [6]: 1 # Limiting current dataset to 10000 rows
 2 df = df[:10000]

In [7]: 1 print('No. of datapoints/rows: {}'.format(df.shape[0]))
 2 print('No. of features/columns: {}'.format(df.shape[1]))

No. of datapoints/rows: 10000
 No. of features/columns: 2

In [8]: 1 print("Feature names: \n{}".format(df.columns))

Feature names:
 Index(['review', 'sentiment'], dtype='object')

```
In [9]: 1 #Check null and missing values
        2 # Calculate the number of missing values in each column using isna()
        3 df.isna().sum()
```

```
Out[9]: review      0
        sentiment    0
        dtype: int64
```

```
In [10]: 1 # Calculate the number of missing values in each column
         2 df.isnull().sum()
```

```
Out[10]: review      0
         sentiment    0
         dtype: int64
```

```
In [11]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   review      10000 non-null   object
1   sentiment    10000 non-null   object
dtypes: object(2)
memory usage: 156.4+ KB
```

```
In [12]: 1 # Check original shape of the dataset
         2 df.shape
```


```
Out[12]: (10000, 2)
```

```
In [13]: 1 df['sentiment'].value_counts()
```

```
Out[13]: positive      5028
         negative      4972
         Name: sentiment, dtype: int64
```

```
In [14]: 1 import seaborn as sns  
2 sns.countplot(x='sentiment',data=df)
```

Out[14]: <Axes: xlabel='sentiment', ylabel='count'>




```
In [15]: 1 positive_review=list(df[df['sentiment']=='positive']['review'][:100])  
2 negative_review=list(df[df['sentiment']=='negative']['review'][:100])
```

```
In [16]: 1 from wordcloud import WordCloud, STOPWORDS
2 from matplotlib import pyplot as plt
3 stopwords=set(STOPWORDS)
4 stopwords
```

```
Out[16]: {'a',
'about',
'above',
'after',
'again',
'against',
'all',
'also',
'am',
'an',
'and',
'any',
'are',
"aren't",
'as',
'at',
'be',
'because',
'been',
...}
```

```
In [17]: 1 def create_cloud(string, title=None):
2     cloud = WordCloud(height=1080,
3                       width=1920,
4                       background_color='white',
5                       min_font_size=10,
6                       stopwords=STOPWORDS).generate(string)
7     plt.figure(figsize=(10,20))
8     plt.imshow(cloud)
9     plt.axis("off")
10    if title:
11        plt.title(title, fontdict={'fontsize':24})
12    plt.show()
```

```
1 create_cloud(' '.join(positive_review).lower(), 'words in positive review')
```



```
1 create_cloud(' '.join(negative_review).lower(), 'words in negative review')
```



```

In [20]: 1 def text_processing(data):
2         from bs4 import BeautifulSoup
3         import re
4         def decontracted(phrase):
5             # specific
6             phrase= re.sub(r'<br /><br />', ' ', phrase)
7             phrase = re.sub(r"won't", "will not", phrase)
8             phrase = re.sub(r"can't", "can not", phrase)
9
10            # general
11            phrase = re.sub(r"n't", " not", phrase)
12            phrase = re.sub(r"\ 're", " are", phrase)
13            phrase = re.sub(r"\ 's", " is", phrase)
14            phrase = re.sub(r"\ 'd", " would", phrase)
15            phrase = re.sub(r"\ 'll", " will", phrase)
16            phrase = re.sub(r"\ 't", " not", phrase)
17            phrase = re.sub(r"\ 've", " have", phrase)
18            phrase = re.sub(r"\ 'm", " am", phrase)
19            phrase = re.sub(r'""', " ", phrase)
20            return phrase
21        stopwords=set(STOPWORDS)
22
23        # Combining all the above sentence
24        from tqdm import tqdm
25        preprocessed_reviews = []
26        # tqdm is for printing the status bar
27        for sentence in tqdm(data['review'].values):
28            sentence = re.sub(r"http\S+", "", sentence)
29            sentence = BeautifulSoup(sentence, 'lxml').get_text()
30            sentence = decontracted(sentence)
31            sentence = re.sub(r"\S*\d\S*", "", sentence).strip()
32            # https://gist.github.com/sebleier/554280
33            sentence = ' '.join(e.lower() for e in sentence.split() if e not
34                                in stopwords)
35            preprocessed_reviews.append(sentence.strip())
36
37        from nltk.stem import PorterStemmer
38
39        porter = PorterStemmer()
40        list_of_sentence=[]
41        for sentence in preprocessed_reviews:
42            words_in_sentence=[]
43            for words in sentence.split():
44                words_in_sentence.append(porter.stem(words))
45
46            list_of_sentence.append(' '.join(words_in_sentence))
47        return(list_of_sentence)

```

```
In [21]: 1 x=text_processing(df[:5000])
```

```
100%|██████████| 5000/5000 [00:02<00:00, 2345.01it/s]
```



```
In [22]: 1 df = df[:5000]
```

```
In [23]: 1 df['cleaned_review']=x
```

```
In [24]: 1 df.head()
```

```
Out[24]:
```

	review	sentiment	cleaned_review
0	One of the other reviewers has mentioned that ...	positive	one review mention watch oz episod will hooked...
1	A wonderful little production. The...	positive	a wonder littl production. the film techniqu u...
2	I thought this was a wonderful way to spend ti...	positive	i thought wonder way spend time hot summer wee...
3	Basically there's a family where a little boy ...	negative	basic famili littl boy (jake) think zombi clos...
4	Petter Mattei's "Love in the Time of Money" is...	positive	petter mattei love time money visual stun film...

```
In [25]: 1 X = df['cleaned_review']
2 Y = df['sentiment']
```

```
In [26]: 1 Y = list(Y)
2 for i in range(len(Y)):
3     if Y[i]=='positive':
4         Y[i]=1
5     else:
6         Y[i]=0
7
8 df['sentiment_score']=Y
9
10 Y=df['sentiment_score']
```

In [27]:

1 df

Out[27]:

	review	sentiment	cleaned_review	sentiment_score
0	One of the other reviewers has mentioned that ...	positive	one review mention watch oz episod will hooked...	1
1	A wonderful little production. The...	positive	a wonder littl production. the film techniqu u...	1
2	I thought this was a wonderful way to spend ti...	positive	i thought wonder way spend time hot summer wee...	1
3	Basically there's a family where a little boy ...	negative	basic famili littl boy (jake) think zombi clos...	0
4	Petter Mattei's "Love in the Time of Money" is...	positive	petter mattei love time money visual stun film...	1
...
4995	An interesting slasher film with multiple susp...	negative	an interest slasher film multipl suspects.incl...	0
4996	i watched this series when it first came out i...	positive	watch seri first came year old watch best frie...	1
4997	Once again Jet Li brings his charismatic prese...	positive	onc jet li bring charismat presenc movi screen...	1
4998	I rented this movie, after hearing Chris Gore ...	negative	i rent movie, hear chri gore say someth effect...	0
4999	This was a big disappointment for me. I think ...	negative	thi big disappoint me. i think worst mastroian...	0

5000 rows × 4 columns

In [28]:

```
1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(
3     X[:5000], Y[:5000], test_size=0.3, random_state=0)
```

In [29]:

```
1 X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[29]: ((3500,), (1500,), (3500,), (1500,))

```
In [30]: 1 X_train
```

```
Out[30]: 2858   thi film littl recommend it, though littl brea...
1559   thi kind inspir saccharin enough make sick. it...
1441   i confess i know involved, i forerunn the plan...
2179   thi movi delic argument search someth make thi...
1390   there dvd publish uk code cover, asin, vfc dis...
...
4931   françoi villon real-lif poet rogu live pari ce...
3264   i want one - situat rich, set unusu interestin...
1653   yeah, sum up. thi movi horrifying. two minut i...
2607   i honest, i realli good time watch she man. de...
2732   ag excel present drama, suspens thriller rare ...
Name: cleaned_review, Length: 3500, dtype: object
```

```
In [31]: 1 list(y_test).count(0)
```

```
Out[31]: 777
```

```
In [32]: 1 from sklearn.feature_extraction.text import CountVectorizer
2
3 vectorizer = CountVectorizer()
4 X_train_bow = vectorizer.fit_transform(X_train)
5 X_test_bow = vectorizer.transform(X_test)
```

```
In [33]: 1 X_train_bow.shape, X_test_bow.shape
```

```
Out[33]: ((3500, 30112), (1500, 30112))
```

```
In [34]: 1 X_train.shape
```

```
Out[34]: (3500,)
```

1. KNN

```

In [35]: 1 from sklearn.neighbors import KNeighborsClassifier
2 from sklearn.metrics import accuracy_score, f1_score
3 for i in range(25,30):
4
5     print('K',i)
6
7     # initialization
8     neigh = KNeighborsClassifier(n_neighbors=i)
9
10    # Training
11    neigh.fit(X_train_bow, y_train)
12
13    # Test the training data
14    y_pred_train = neigh.predict(X_train_bow)
15    accuracy_train = accuracy_score(y_pred_train,y_train)
16    f1_train = f1_score(y_pred_train,y_train)
17
18    # Test the test data
19    y_pred_test = neigh.predict(X_test_bow)
20    accuracy_test = accuracy_score(y_pred_test,y_test)
21    f1_test = f1_score(y_pred_test,y_test)
22
23    print(accuracy_train,accuracy_test)
24    print(f1_train,f1_test)
25    print()

```

```

K 25
0.7245714285714285 0.6406666666666667
0.7169700528479155 0.6254343293954134

```

```

K 26
0.7102857142857143 0.634
0.6805293005671077 0.5887640449438202

```

```

K 27
0.7222857142857143 0.644
0.7136122569239836 0.6307053941908713

```

```

K 28
0.7077142857142857 0.6493333333333333
0.680811232449298 0.6120943952802359

```

```

K 29
0.7131428571428572 0.6506666666666666
0.7079697498545666 0.6401098901098901

```

```

In [ ]: 1 from sklearn.neighbors import KNeighborsClassifier
2 from sklearn.metrics import accuracy_score
3 from sklearn.metrics import f1_score
4
5 # initialization
6 neigh = KNeighborsClassifier(n_neighbors=28)
7
8 # Training
9 neigh.fit(X_train_bow, y_train)
10
11 # Test the training data
12 y_pred_train = neigh.predict(X_train_bow)
13 accuracy_train = accuracy_score(y_pred_train,y_train)
14 f1_train = f1_score(y_pred_train,y_train)
15
16
17 # Test the test data
18 y_pred_test = neigh.predict(X_test_bow)
19 accuracy_test = accuracy_score(y_pred_test,y_test)
20 f1_test = f1_score(y_pred_test,y_test)
21
22
23 print(accuracy_train,accuracy_test)
24 print(f1_test,f1_test)

```

```

In [36]: 1 from sklearn.metrics import classification_report
2 target_names = ['Positive', 'Negative']
3 print(classification_report(y_pred_test, y_test, target_names=target_names))
4 print(classification_report(y_pred_train, y_train, target_names=target_names))

```

	precision	recall	f1-score	support
Postive	0.66	0.66	0.66	767
Negative	0.64	0.64	0.64	733
accuracy			0.65	1500
macro avg	0.65	0.65	0.65	1500
weighted avg	0.65	0.65	0.65	1500


	precision	recall	f1-score	support
Postive	0.73	0.71	0.72	1807
Negative	0.70	0.72	0.71	1693
accuracy			0.71	3500
macro avg	0.71	0.71	0.71	3500
weighted avg	0.71	0.71	0.71	3500

```
In [37]: 1 from sklearn.metrics import confusion_matrix
2 c= confusion_matrix(y_pred_train, y_train)
3 c
```

```
Out[37]: array([[1279,  528],
 [ 476, 1217]], dtype=int64)
```

```
In [38]: 1 import seaborn as sns
2 sns.heatmap(c, annot=True)
```

```
Out[38]: <Axes: >
```



```
In [39]: 1 from sklearn.model_selection import GridSearchCV
2
3 parameters = {'n_neighbors':list(range(10,30,2))}
4 neigh = KNeighborsClassifier()
5
6 clf = GridSearchCV(neigh, parameters)
7 clf.fit(X_train_bow, y_train)
```

```
Out[39]: GridSearchCV(estimator=KNeighborsClassifier(),
                      param_grid={'n_neighbors': [10, 12, 14, 16, 18, 20, 22, 24, 26,
                      28]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [40]: 1 clf.best_params_

Out[40]: {'n_neighbors': 28}

In [41]:

```

1 neigh = KNeighborsClassifier(n_neighbors=10, p=2)
2 neigh.fit(X_train_bow, y_train)
3
4 y_pred_train = clf.predict(X_train_bow)
5 f1_train = f1_score(y_pred_train, y_train)
6 print(f1_train)
7 print(classification_report(y_pred_train, y_train, target_names=target_name))

```

0.680811232449298

	precision	recall	f1-score	support
Postive	0.79	0.68	0.73	2040
Negative	0.63	0.75	0.68	1460
accuracy			0.71	3500
macro avg	0.71	0.71	0.71	3500
weighted avg	0.72	0.71	0.71	3500

In [42]:

```

1 y_pred_test = clf.predict(X_test_bow)
2 f1_test = f1_score(y_pred_test, y_test)
3 print(f1_test)
4 print(classification_report(y_pred_test, y_test, target_names=target_name))

```

0.6120943952802359

	precision	recall	f1-score	support
Postive	0.72	0.64	0.68	867
Negative	0.57	0.66	0.61	633
accuracy			0.65	1500
macro avg	0.65	0.65	0.65	1500
weighted avg	0.66	0.65	0.65	1500

2. Decision Tree

In [43]:

```

1 from sklearn.tree import DecisionTreeClassifier
2 clf = DecisionTreeClassifier()
3 clf.fit(X_train_bow, y_train)
4 DecisionTreeClassifier()
5 y_pred = clf.predict(X_test_bow)
6 accuracy = metrics.accuracy_score(y_test, y_pred)
7 accuracy

```

Out[43]: 0.682

3. Random Forest

```
In [44]: 1 from sklearn.ensemble import RandomForestClassifier
2 rf_classifier = RandomForestClassifier(n_estimators=100,
3 random_state=42)
4 rf_classifier.fit(X_train_bow, y_train)
5 RandomForestClassifier(random_state=42)
6 y_pred = rf_classifier.predict(X_test_bow)
7 accuracy = metrics.accuracy_score(y_test, y_pred)
8 accuracy
```

Out[44]: 0.828

4. Boosting Algorithm

```
In [45]: 1 from sklearn.ensemble import AdaBoostClassifier
2 base_classifier = DecisionTreeClassifier(max_depth=3)
3 adaboost_classifier = AdaBoostClassifier(base_classifier,
4 n_estimators=50, random_state=42)
5 adaboost_classifier.fit(X_train_bow, y_train)
6 AdaBoostClassifier(estimator=DecisionTreeClassifier(max_depth=1),
7 random_state=42)
8 y_pred = adaboost_classifier.predict(X_test_bow)
9 accuracy = metrics.accuracy_score(y_test, y_pred)
10 accuracy
```

Out[45]: 0.7646666666666667

Boosting Algorithm is the best accuracy

```
In [47]: 1 accuracy
```

Out[47]: 0.7646666666666667

```
In [ ]: 1
```

5. Naive Bayes classifier

```
In [49]: 1 from sklearn.naive_bayes import GaussianNB
2 from sklearn.naive_bayes import BernoulliNB
3 from sklearn.naive_bayes import MultinomialNB
4 from sklearn import metrics
```



```
In [59]: 1 # BernoulliNB
2 B_classifier = BernoulliNB()
3 B_classifier.fit(X_train_bow, y_train)
4 BernoulliNB()
5 predictions_B = B_classifier.predict(X_test_bow)
6 accuracy_B = metrics.accuracy_score(y_test, predictions_B)
7 accuracy_B
```

Out[59]: 0.816

6. Logistic Regression

```
In [46]: 1 from sklearn import linear_model
2 lrg = linear_model.LogisticRegression()
3 lrg.fit(X_train_bow, y_train)
4 LogisticRegression()
5 lrg.score(X_test_bow, y_test)
```

Out[46]: 0.8513333333333334

```
In [ ]: 1
```